Sign Tests: Understanding the Direction of Change

The Definition

Welcome to the study of Sign Tests. This is a 'nonparametric' statistical test. Nonparametric is just a fancy word for a **flexible method that does not require your data to follow a perfect,** symmetrical bell-shaped curve (normal distribution). It is a tool used when our data is a bit **messy,** or the sample size is too small for standard tests like the t-test.

The Focus

Unlike other tests that obsess over exact numbers and averages, the Sign Test cares primarily about the *direction* of the difference. It asks a simple question: "Is the value higher (+) or lower (-)?" rather than "How much higher is it?" It strips away the magnitude to focus on the trend.

The Goal

There are two main ways we use this tool. We can test a single group against a specific guess (called a Hypothesized Median), or we can compare two related groups—like a "before and after" scenario—to see if there has been a significant shift in performance or opinion.

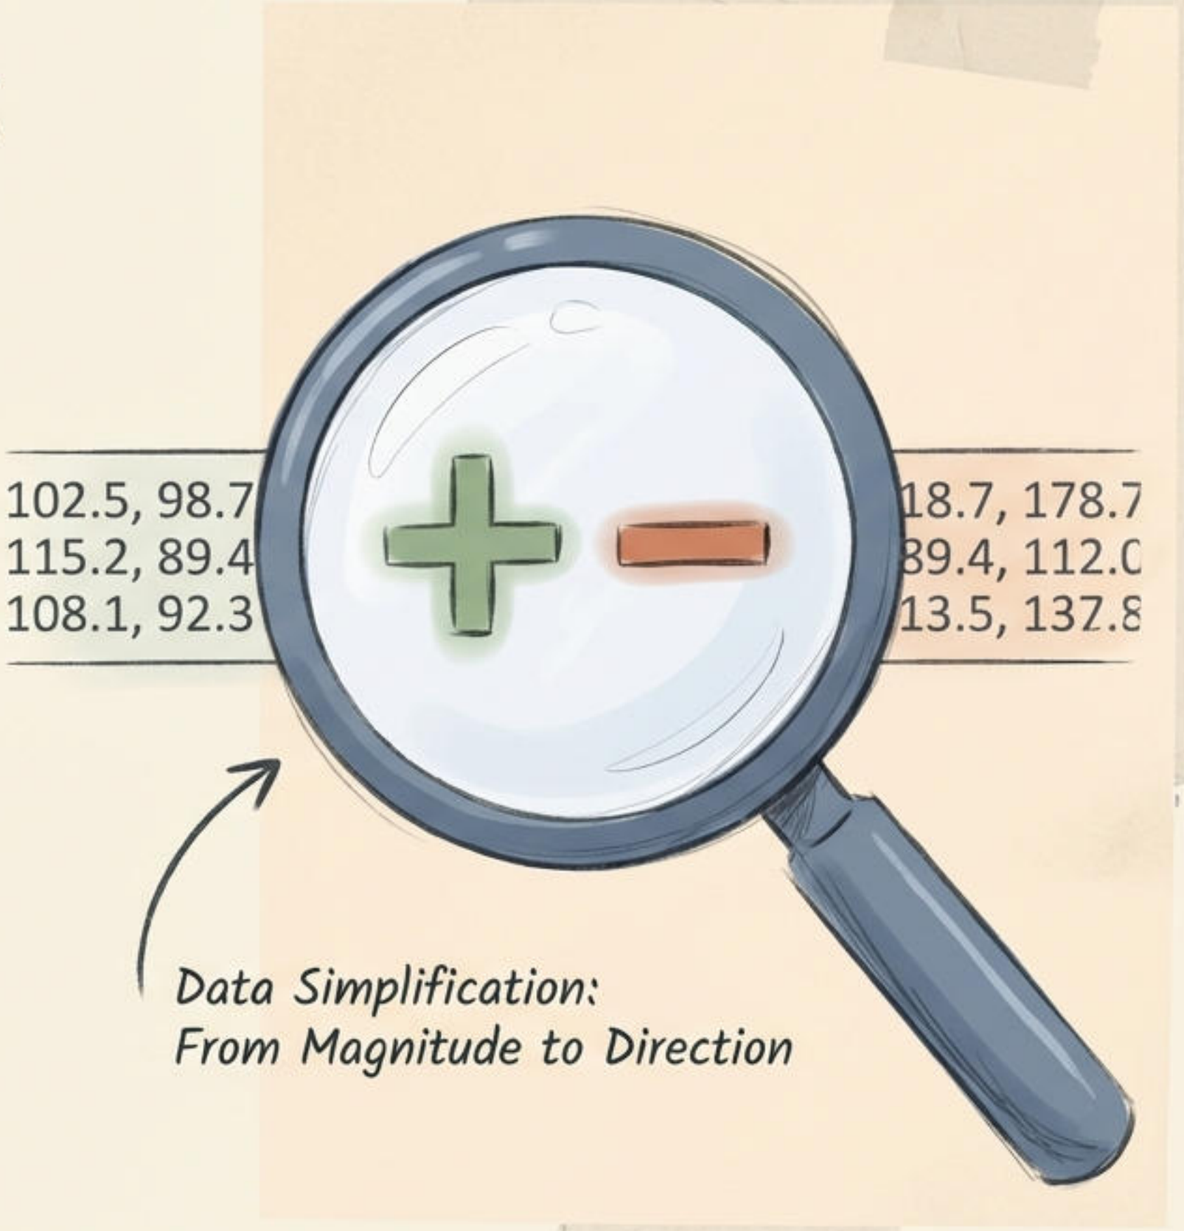

NotebookLM

Why We Focus on the Median, Not the Mean

The Problem with Averages

In standard statistics, we often look at the arithmetic mean (the average). However, the mean is easily ruined by "outliers"—extreme values that pull the average up or down incorrectly. Imagine a room of poor students; if a billionaire walks in, the "average" wealth skyrockets, but the students are still poor.

The Median Solution

The Sign Test focuses on the "location-wise central point," known as the median. The median is simply the middle value where half the group is higher and half is lower. As you can see in the illustration, the median is much more stable when the data contains extreme values that might skew the result.

When to Use It

We use this test specifically when we cannot assume the population is "normal" (symmetrical). By ignoring the magnitude (size) of the numbers and only looking at their position relative to the median, we get a safer, more reliable result for difficult or skewed data sets.

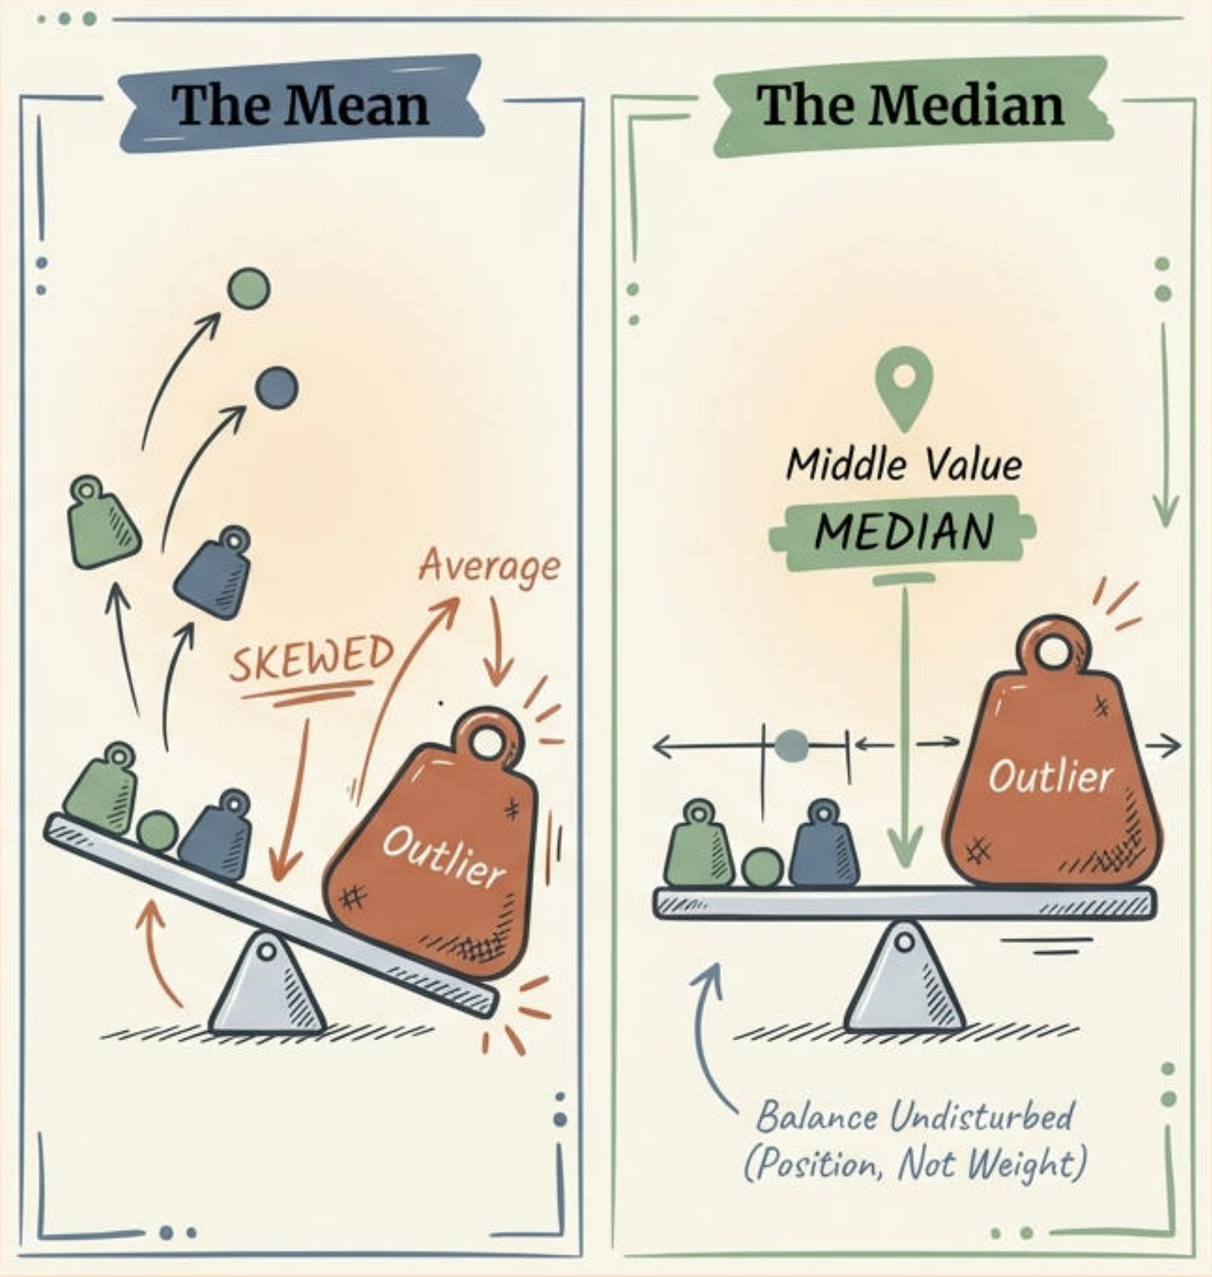

NotebookLM

The One-Sample Test: Assigning the Signs

The Setup

In the "One-Sample" scenario, we have a single group of random data, and we want to compare it to a specific value we guess is true.

M_0

We call this guess the "Hypothesized Median" (M_0). For example, guessing that the median age of a group is 30.

The Comparison

The mechanical process is straightforward. We take every single number in our collected data and compare it to our guess (M_0).

We don't care if the number is slightly bigger or massively bigger; we prry or massively bigger; we only care about the relationship.

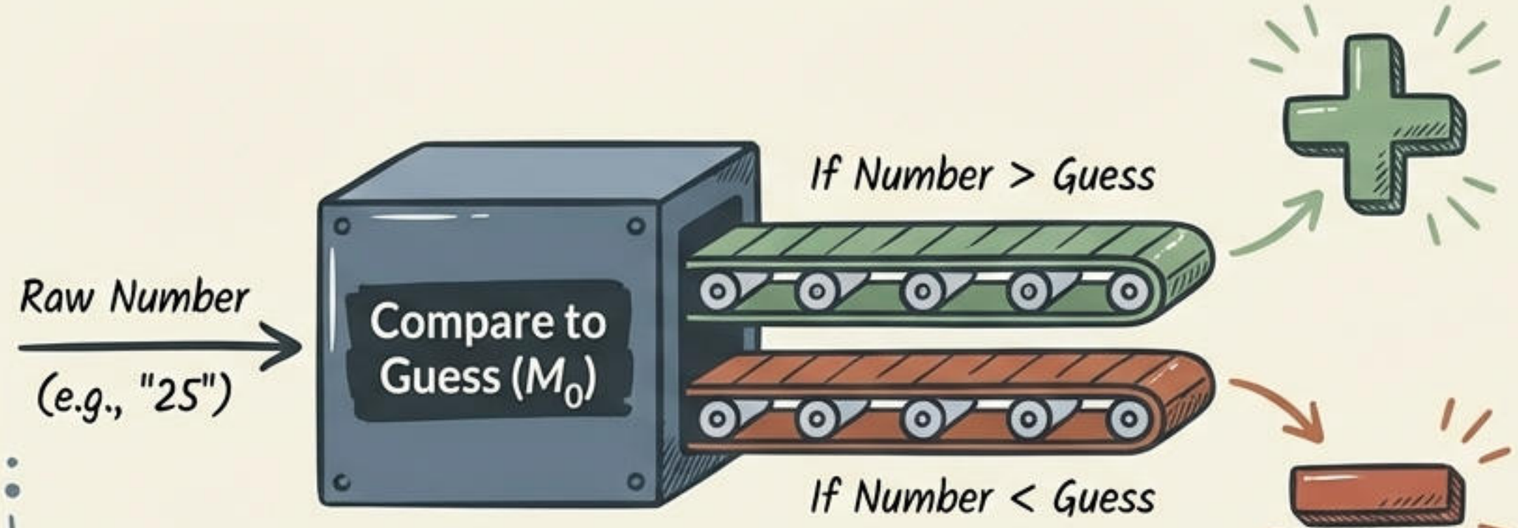

The Transformation

If a data point is greater than our guess ($X > M_0$), we replace the number with a Plus sign (+).

If the data point is smaller than our guess ($X < M_0$), we replace it with a Minus sign (-).

This transforms a complex list of numbers into a simple string of signs, ready for analysis.

NotebookLM

Handling Zeroes and the Null Hypothesis

The Rule of Zeroes.

What happens if a data point is exactly equal to our guess (M_0)? In this case, the difference is zero. The rule is strictly to *discard* these observations. We act as if they never happened. This is crucial because it reduces our total sample size (denoted as n).

The Null Hypothesis (H_0).

In statistics, the Null Hypothesis (H_0) is our "innocent until proven guilty" stance. We assume there is *no* difference and that our guess is correct. We maintain this assumption unless the data screams otherwise.

The Logic of Probability.

If H_0 is true (our guess is right), the data should be randomly scattered around the median. This means we should expect roughly 50% plus signs and 50% minus signs, just like a balanced coin flip. If we see a massive imbalance (like 90% plus signs), we know our original guess was probably wrong.

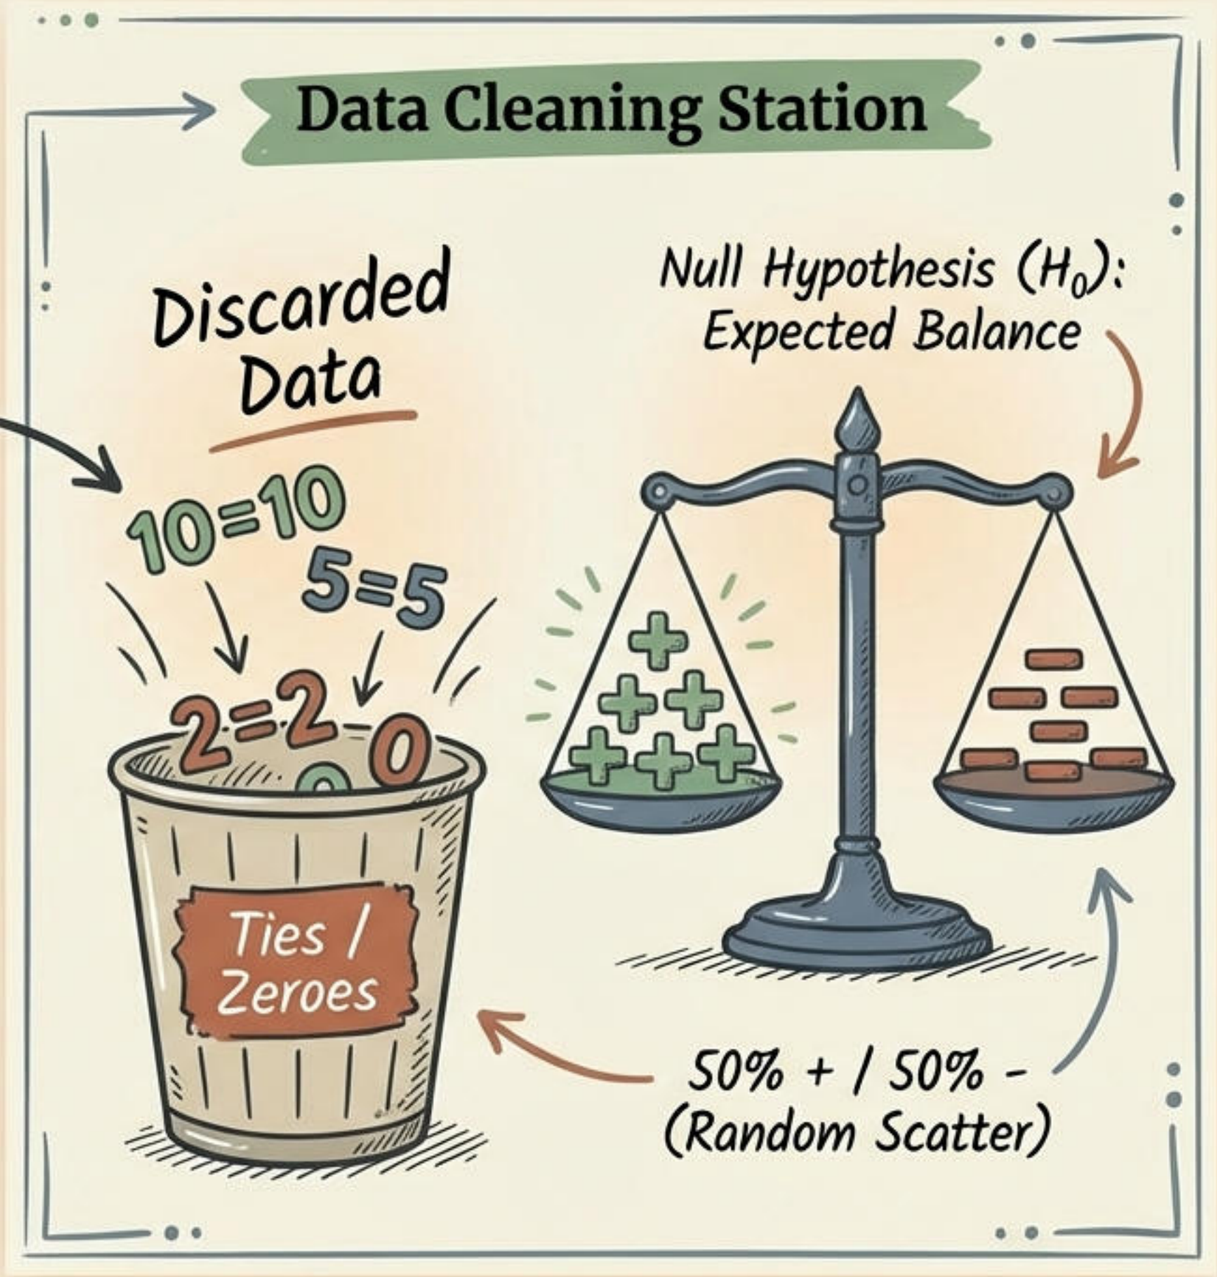

NotebookLM

The Test Statistic (S) and Binomial Logic

Defining S

We need a number to measure the result. We call this the test statistic, "S". In a One-Sample Sign Test, S is simply the count of the sign that appears most frequently. If we have 2 pluses and 10 minuses, S is 10.

The Coin Flip Analogy

Think of this like flipping a coin. Under normal conditions (a fair coin), getting Heads or Tails is a 50/50 shot. If you flipped a coin 10 times and got Heads every single time, you would suspect the coin is rigged. The Sign Test uses the exact same logic.

The Binomial Distribution

Statisticians use the "Binomial Distribution" (a formula for yes/no outcomes) to calculate exactly *how* rare a specific pattern is. We use this math to find the "p-value," which tells us the probability of seeing our specific count of signs just by luck.

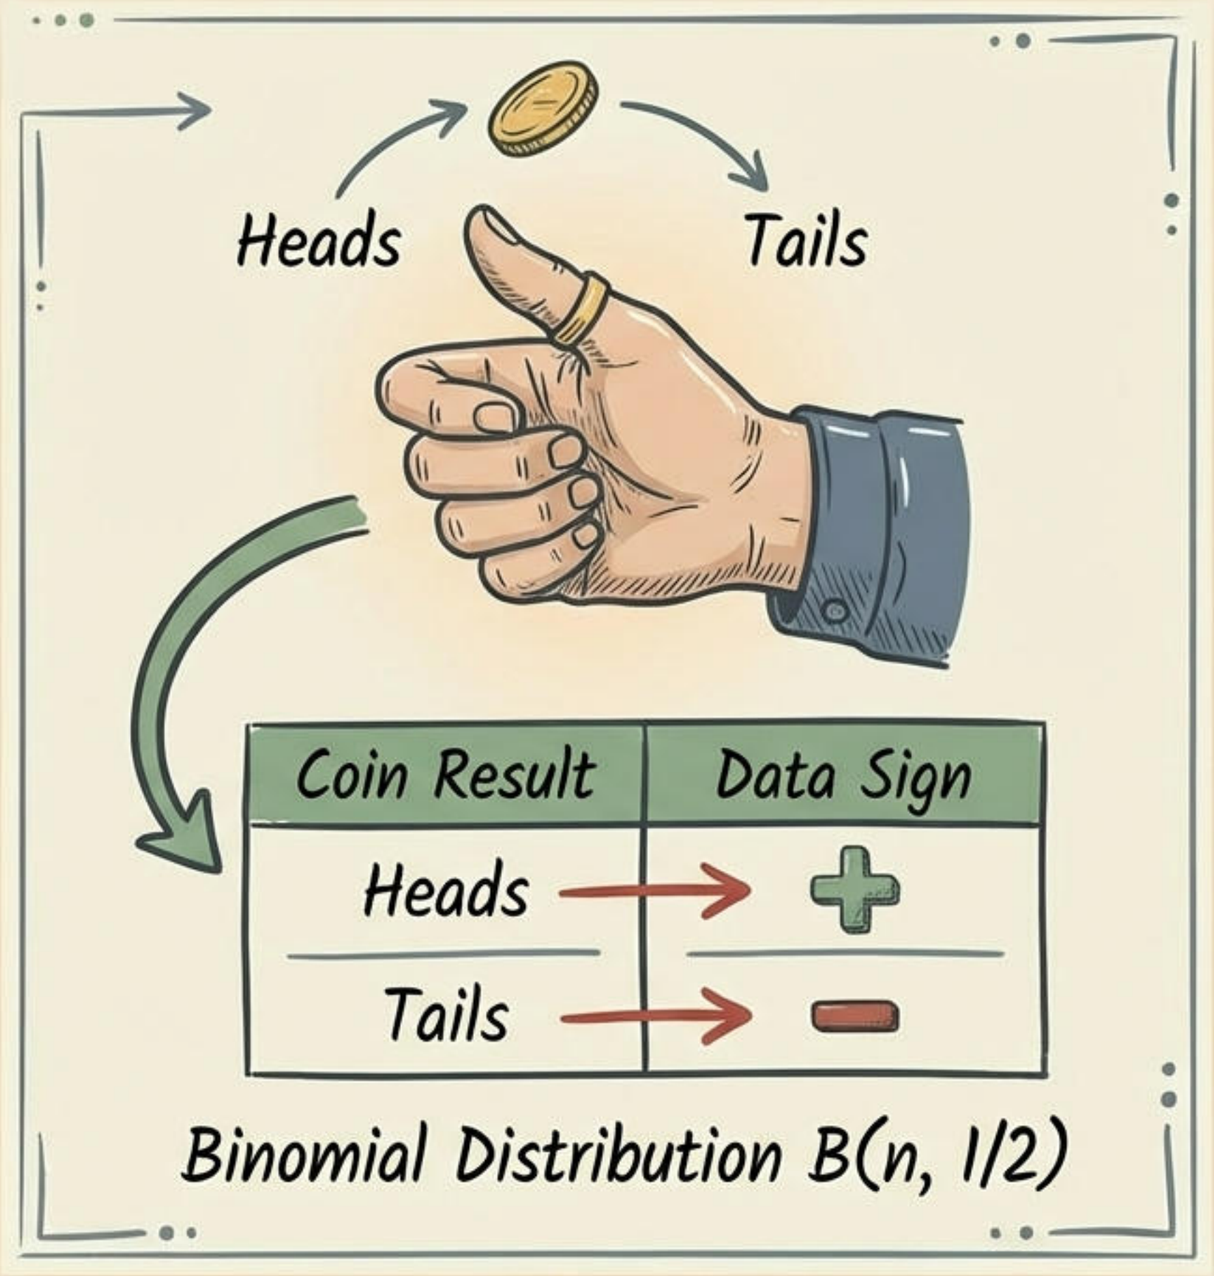

NotebookLM

Case Study: Professional Golfers (The Setup)

The Scenario

Let's apply this to a real story. We have data from 11 professional golfers playing 4 rounds at a city club. We want to know if these professionals are performing significantly better than the course average.

The Hypothesis

The course average is 284. Our **Null Hypothesis (H_0)** is that the median score of these golfers is **exactly 284** (they are average). Our **Alternative Hypothesis (H_1)** is that their median is *less* than 284. In golf, a lower score is better, so we are testing if they are skilled players.

The Data

We collect the total scores for the 11 players. The raw numbers are: 280, 282, 292, 273, 283, 275, 284, 282, 279, and 281. The list looks random, but the Sign Test will help us see the pattern hidden within.

NotebookLM

Processing the Golf Data

Assigning Signs

We compare each golfer's score to the course average of 284. For example, the first player scored 280. Since 280 is lower than 284, they get a Minus (–). The third player scored 292, which is higher, so they get a Plus (+).

Discarding the Tie

Notice the player who scored exactly 284. Because this score equals our guess, it provides no information about direction. We discard this result entirely. We started with 11 golfers, but now our sample size (n) is reduced to 10.

Calculating S

Now we count. We are left with exactly 1 plus sign (+) and 9 minus signs (–). Since we are testing if they are better (lower score), or simply looking for the most frequent sign, we focus on the minuses. Our test statistic S is 9. This is a very lopsided result.

Equal to Guess

Raw Data	280	282	292	273	283	275	~~284~~	282	279	281
Comparison	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
Signs	–	–	+	–	–	–	~~DISCARD~~	–	–	–

$S = 9$ (Count of Minuses)

NotebookLM

The Verdict: Are the Golfers Better?

Finding the P-Value

Using a Binomial reference table, we look up the probability of getting 9 or more minuses out of 10 tries just by random chance. The math tells us this probability (p-value) is 0.0108.

Understanding the Probability

What does 0.0108 mean? It means there is essentially a 1% chance that these golfers are just average players having a lucky day. It is an extremely rare event under the assumption that they are "normal" players.

The Decision

In science, we usually set a "safety limit" (significance level) of 0.05 (5%). If the chance of luck is lower than 5%, we accept the result as real. Since 1% is much smaller than 5%, the result is **statistically significant**. We reject the Null Hypothesis. The professional golfers truly average less than 284.

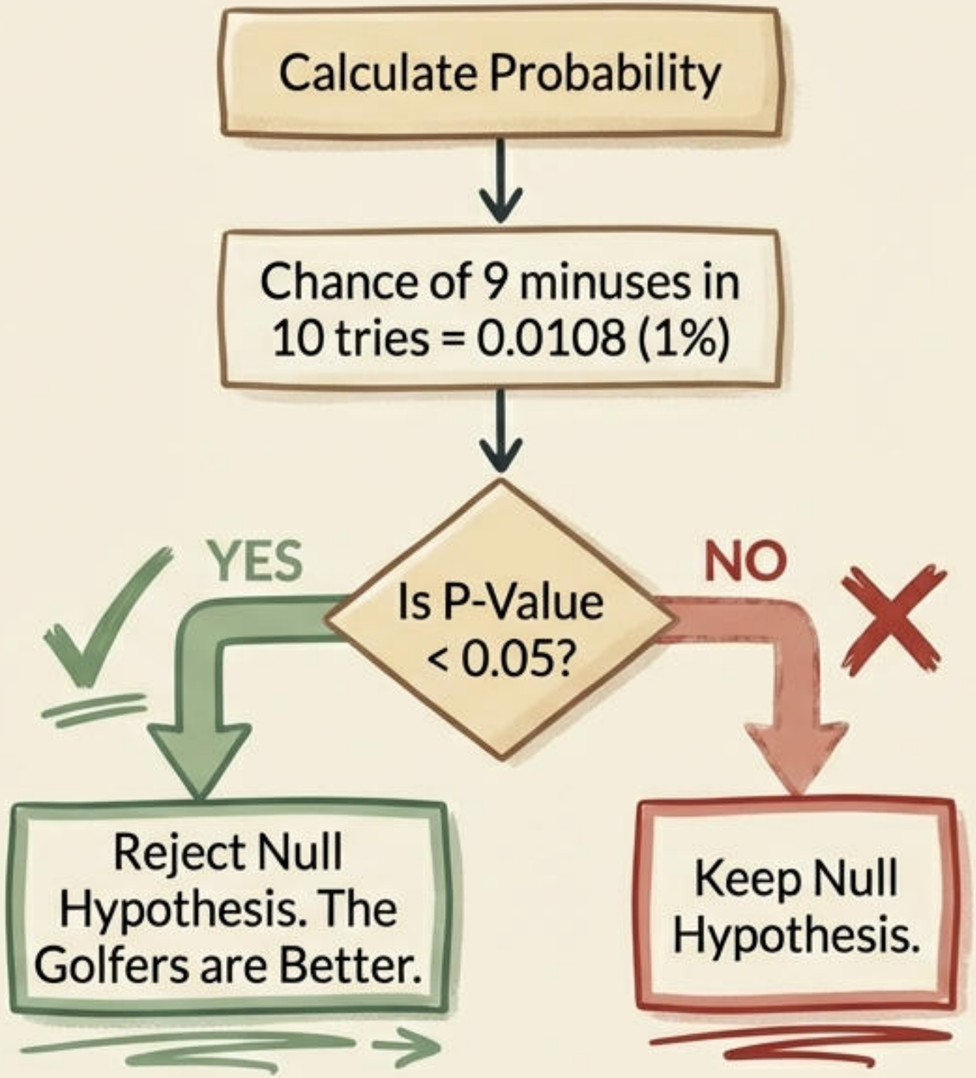

NotebookLM

The Two-Sample (Paired) Sign Test

When to Use It

The second major way we use the Sign Test is for 'paired' observations. This is incredibly common in medical or psychological studies. Think of 'Before and After' treatment (the same person tested twice) or matched pairs (like comparing a husband's opinion vs. his wife's).

The Mechanics

In this version, we are no longer comparing a number to a fixed guess (M_0). Instead, we are looking at the relationship *between* the pairs. We want to know: does the second group consistently score higher or lower than the first?

The Procedure

The math remains simple. For every pair, we subtract the second score from the first (Score A minus Score B). If the result is positive, we mark a (+). If negative, we mark a (-). If the scores are identical, the difference is zero, and we discard the pair.

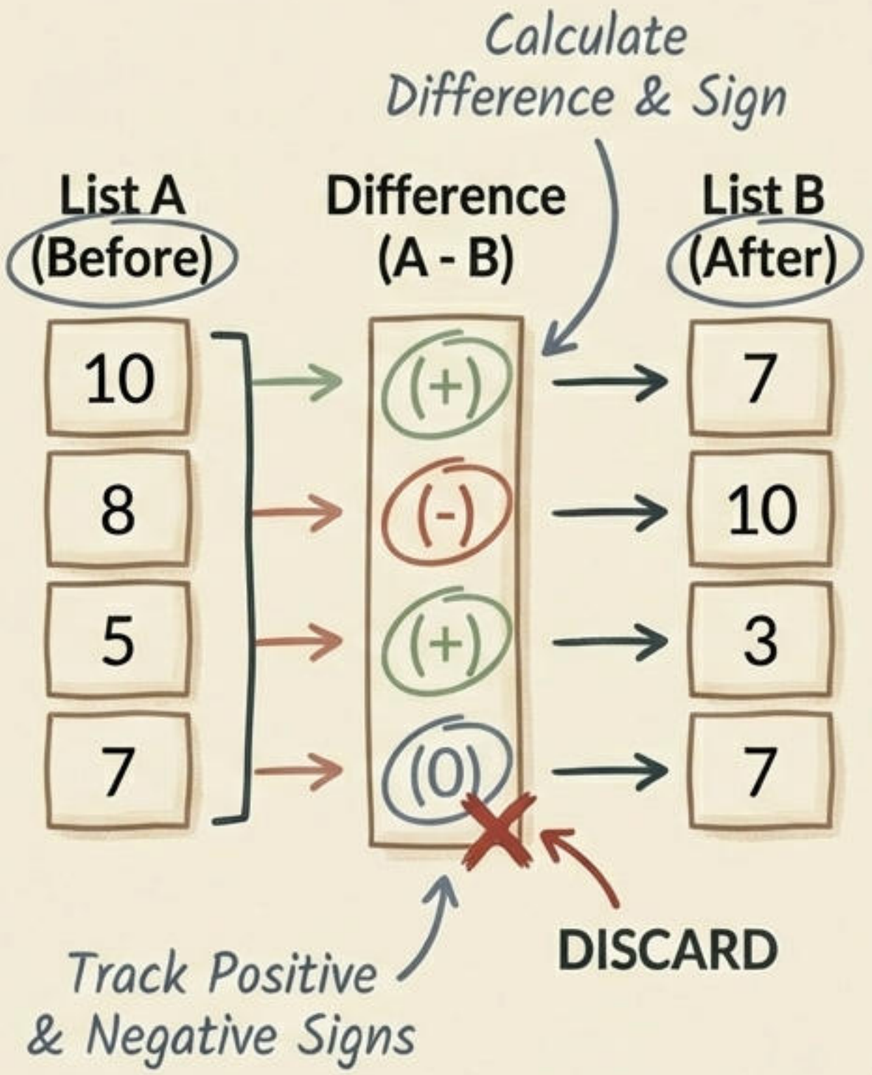

NotebookLM

Case Study: The Archaeologists (Setup)

The Scenario: Two archaeologists, Person A and Person B, have been digging for artifacts at an ancient cliff dwelling for 30 days. We want to know if Archaeologist A is more skilled at finding artifacts than Archaeologist B.

The Data: We have a detailed record of their performance. For each of the 30 days, we recorded exactly how many items each person found. On Day 1, A found 1 and B found 0. On Day 2, they both found 0.

The Hypothesis: Our Null Hypothesis (H_0) assumes they are equally skilled ($Median_A = Median_B$). Our Alternative Hypothesis (H_1) is that A is better than B ($Median_A > Median_B$). We will use the paired sign test to see if A wins significantly more often than B.

Daily Artifact Finds (30 Days)

Arch A	1	0	2	3	1	0	2	2	3	0	1	1	4	1	2	1	3	5	2	1	2	4	1	3	2	0	2	4	2
Arch B	0	0	1	0	2	0	0	1	1	2	0	1	2	1	1	0	2	2	6	0	2	3	0	2	1	0	1	1	0

NotebookLM

Analyzing the Dig: Subtraction and Signs

Calculating Differences

We walk through the 30 days one by one. On Day 1, A found 1 and B found 0. The difference (1 – 0) is positive, so we assign a Plus (+). On Day 3, A found 2 and B found 1. Again, A – B is positive, so another Plus (+).

Handling Ties

On some days, they performed exactly the same (e.g., Day 2: both found 0). Just like in the golf example, these "zero differences" give us no direction. We must discard these 4 tie days from our analysis.

The Resulting Count

We started with 30 days. After throwing out the 4 ties, our **valid sample size** is reduced to $n = 26$. Counting the remaining signs, we find that Archaeologist A won 20 times (+) and lost 6 times (-).

A	B	Sign of (A-B)	
1	0	+	
0	0	0	DISCARD
2	1	+	
3	0	+	
1	2	–	
0	0	0	DISCARD
2	0	+	
2	1	+	
3	1	+	
0	2	–	

Total (+): 20
Total (–): 6
Discarded: 4

NotebookLM

The Problem with Large Samples

Sample Size Limits.

In the golf example, we had 10 data points. It was easy to look up the probability in a small table. However, in the archaeology example, our sample size (n) is 26. When n gets large (usually over 25 or 30), the standard binomial tables become huge, clumsy, and hard to find.

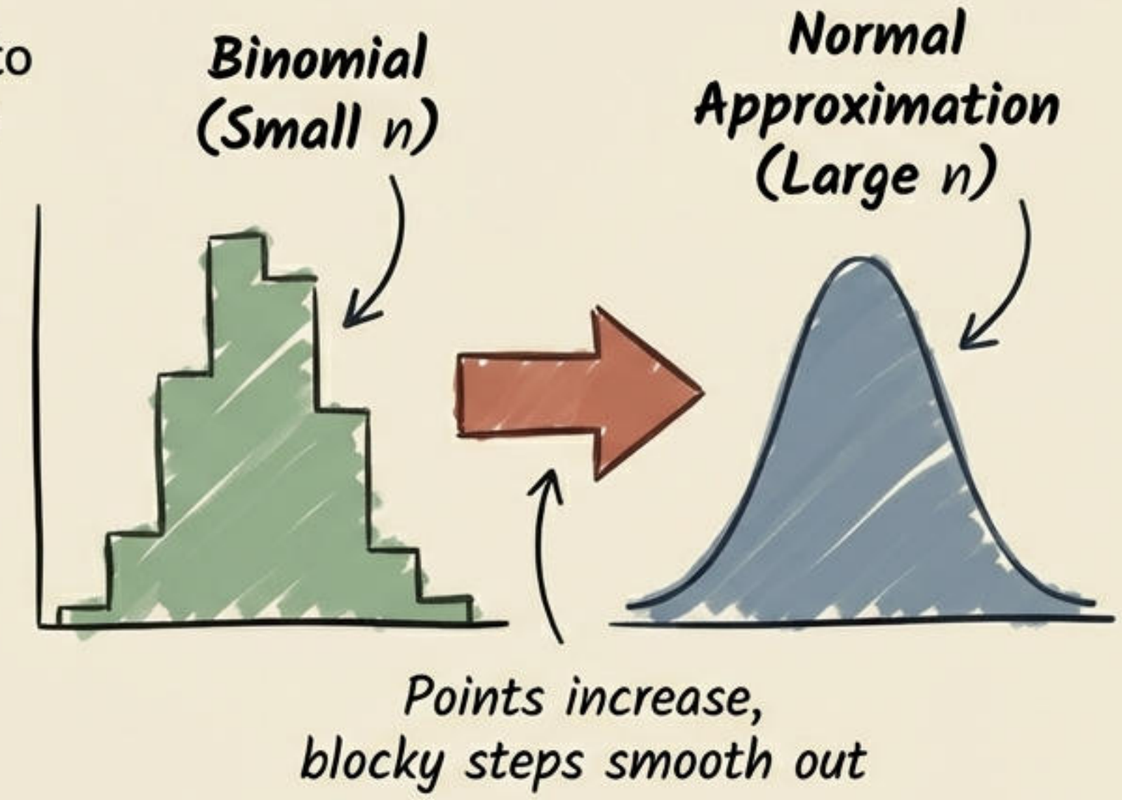

The Normal Approximation.

Fortunately, mathematics gives us a shortcut. As the sample size grows, the Binomial distribution starts starts to look exactly like the famous "Normal Distribution" (the bell curve).

The Z-Score.

Because of this shape-shifting, we can switch methods. Instead of looking up '$' directly, we convert our data into a 'Z-score' using a formula. This standardizes our result so we can use the common Normal Table to find our p-value.

NotebookLM

Calculating the Z-Score

The Formula

To find the Z-score, we use the formula shown below. It might look intimidating, but it is just comparing what we found (S) to what we expected ($n/2$), adjusted for the spread of the data.

Breaking it Down

The term ($n/2$) represents the average we would expect if there was no difference (13 wins for A, 13 wins for B). We subtract this from our actual count ($S = 20$). The bottom part of the fraction ($\sqrt{n/4}$) is the "standard error," which adjusts for the natural variance in the data.

Applying to Archaeology

We simply plug in our numbers. S is 20 (A's wins). n is 26 (total valid days). When we run this calculation, we get a Z-score that tells us exactly how many "standard deviations" away from the average our result falls.

Applying to Archaeology

We simply plug in our numbers. S is 20 (A's wins). n is 26 (total valid days). When we run this calculation, we get a Z-score that tells us **exactly how many "standard deviations" away from the average** our result falls.

Observed Count (20)

Sample Size (26)

$$Z = \frac{S - 0.5 - \frac{n}{2}}{\sqrt{\frac{n}{4}}}$$

Sample Size (26)

Expected Average (13)

NotebookLM

Final Decision: Archaeologist A vs. B

The Result

The formula gives us a Z-score of approximately 2.55. When we check this Z-score against a standard Normal table, it converts to a p-value of 0.0046.

The Meaning

A p-value of 0.0046 is tiny. It means there is less than a 0.5% (half of one percent) chance that Archaeologist A simply got lucky and found more artifacts by accident. This is an extremely rare occurrence.

The Conclusion

Since 0.0046 is significantly lower than our standard safety limit of 0.05, we Reject the Null Hypothesis. We have strong statistical proof that Archaeologist A is indeed better at finding artifacts than Archaeologist B.

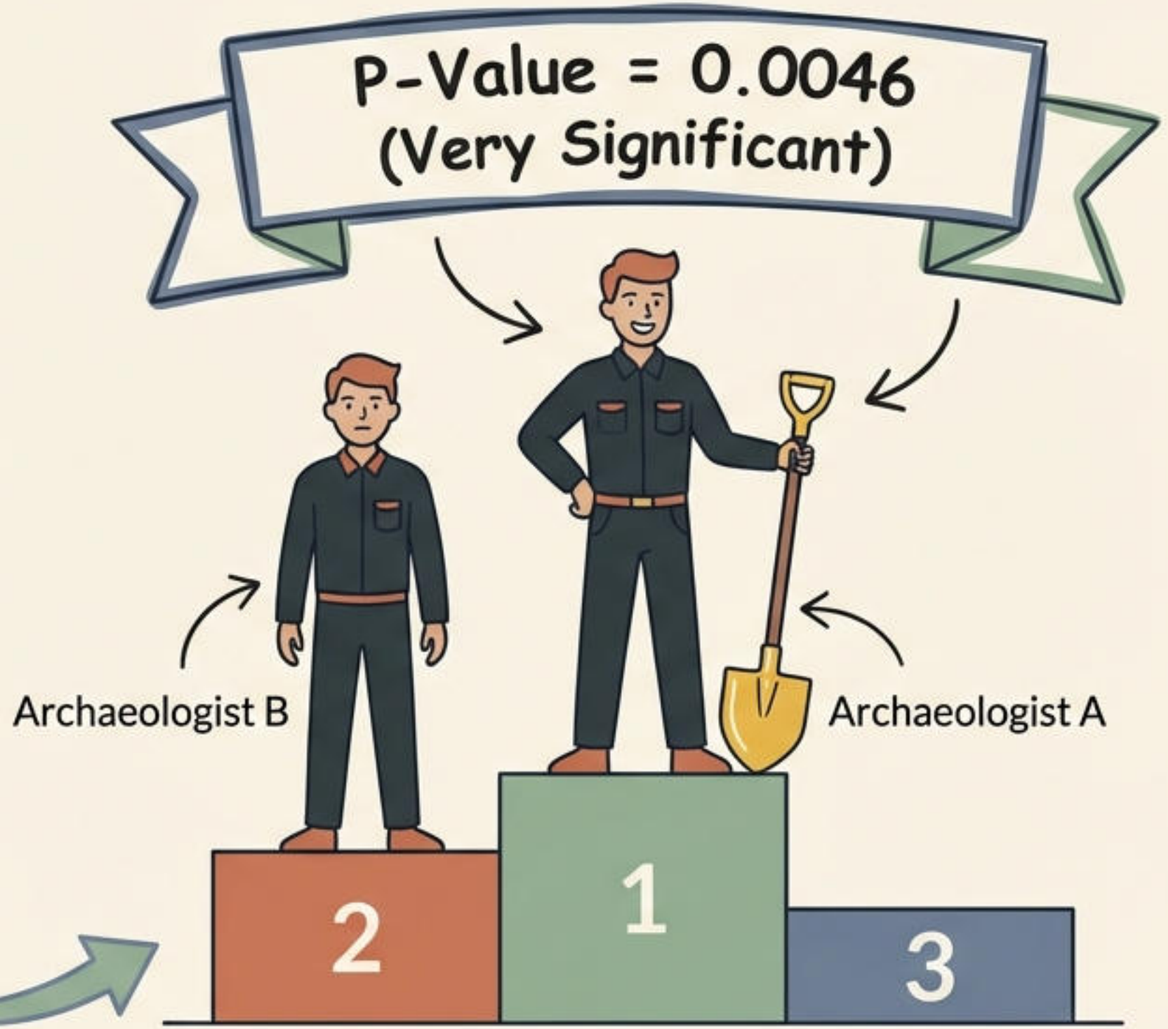

NotebookLM

Summary: The Sign Test Checklist

Step 1: Assign Signs

This is the heart of the test. Whether you are looking at one group (comparing to a guess) or two groups (comparing A vs B), your first job is to convert complex numbers into simple Pluses and Minuses.

Step 2: Clean Data

Never forget to filter. Identify any ties or zeroes—instances where the data equals the guess or the pairs are identical—and remove them from the dataset entirely. Remember to update your sample size (n).

Step 3: Choose Your Path

Look at your new n. If it is small (typically under 25), use the Binomial Table to find your probability. If n is large, use the Z-score formula and the Normal Approximation.

Step 4: Decide

Finally, compare your calculated p-value to the significance level (usually 0.05). If p is low, the Null must go (Reject H_0). If p is high, the difference is likely just luck.

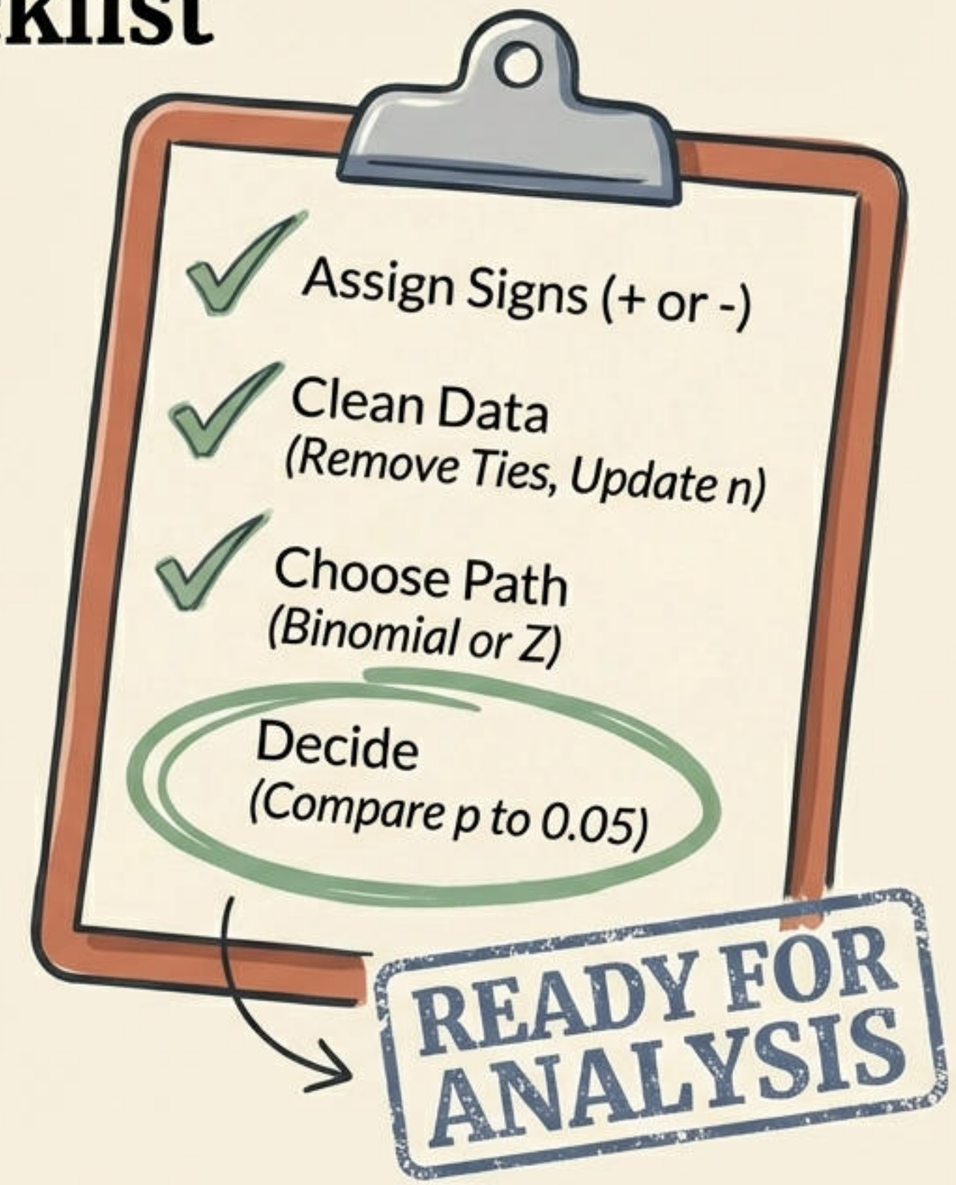

NotebookLM